

training!"

3. If you add any additional variables, make sure you use the same names as the ones used in the class

I've started the code for you below -- how would you finish it?

```
In [1]: import tensorflow as tf
        from os import path, getcwd, chdir

        # DO NOT CHANGE THE LINE BELOW. If you are developing in a local
        # environment, then grab mnist.npz from the Coursera Jupyter Notebook
        # and place it inside a local folder and edit the path to that location
        path = f"{getcwd()}/../tmp2/mnist.npz"

In [2]: # GRADED FUNCTION: train_mnist
        def train_mnist():
            # Please write your code only where you are indicated.
            # please do not remove # model fitting inline comments.

            # YOUR CODE SHOULD START HERE
            class Callback(tf.keras.callbacks.Callback):
                def on_epoch_end(self, epoch, logs={}):
                    if(logs.get('acc')>0.99):
                        print("\nReached 99% accuracy so cancelling training!")
                        self.model.stop_training = True
            # YOUR CODE SHOULD END HERE

            mnist = tf.keras.datasets.mnist

            (x_train, y_train),(x_test, y_test) = mnist.load_data(path=path)
            # YOUR CODE SHOULD START HERE
            x_train = x_train/255
            x_test = x_test/255
            callback = Callback()
            # YOUR CODE SHOULD END HERE
            model = tf.keras.models.Sequential([
                # YOUR CODE SHOULD START HERE
                tf.keras.layers.Flatten(),
                tf.keras.layers.Dense(512, activation=tf.nn.relu),
                tf.keras.layers.Dense(10, activation=tf.nn.softmax)
                # YOUR CODE SHOULD END HERE
            ])

            model.compile(optimizer='adam',
                          loss='sparse_categorical_crossentropy',
                          metrics=['accuracy'])

            # model fitting
            history = model.fit(# YOUR CODE SHOULD START HERE
                               x_train, y_train, epochs = 10, callbacks = [callback]
                               # YOUR CODE SHOULD END HERE
            )
            # model fitting
            return history.epoch, history.history['acc'][-1]
```

In [3]: train\_mnist()